

CS181 / CSCI E-181 Spring 2014 Final Project

David Wihl
davidwihl@gmail.com

Zachary Hendlin
zgh@mit.edu

May 4, 2014

Introduction

To gather sufficient data, we allowed the SampleAgent (provided initially in the code) to play 30 megabytes of ghost training data of games, to get data on feature vectors and associated point values of each.

Classification of Ghosts

We sought to classify ghosts as [0, 1, 2, 3, 5], where all ghosts in category 5 are dangerous (e.g. induce a reward of -1000 points unless a helpful capsule is consumed first).

We explored two methods for classifying the ghosts on the basis of their features.

First, we explored linear support vector machines using SK-Learn's Stochastic Gradient Descent classifier. For classifying the category 5 ghosts, this approach was accurate 90.62 percent of the time. Our analysis found that while differences in the rewards associated with eating ghosts not from class 5 did differ by class, the most important thing for us to measure our performance on is the correct classification of dangerous (class 5 ghosts).

igraph here

We also used logistic regression classification, and achieved somewhat better results, with correct classification of class 5 ghosts 93.25 percent of the time.

We elected to use the logistic regression classification results to predict which class each ghost is in during runtime.

Classification of Capsules and Placebos

We want to determine which of the pills are helpful capsules and which are placebos. To do this, we first plotted the helpful capsules which we collected using the 'd' data collection function.

Here we determined that capsule feature values were very much clustered into three distinct clusters as shown below. `jshow graph`

This suggested that we could either fit Gaussians to these points as part of a generative approach, or apply something like K-means to find the clusters. Because the data were in three dimensions, we opted to use k-means for clustering, and initiated the model with `kmeans++`.

Because the three clusters were all positively identified, simply assigning a new value at run-time (for a capsule which may be a 'good' or 'placebo'), we instead used the 'score' attribute to evaluate the distance from the centroid the data point would have been assigned to. We are agnostic to which centroid the capsule would be assigned to but we are very sensitive to the score assigned (e.g. the distance from the centroid assignment), with larger distances being worse.

From our positive training data, we found values of 0 to -118 as the 'good' range of objective function values. Any value `<-118` would likely be a placebo capsule.

So at runtime we score each capsule based on its 3x1 feature vector and we consider the 'best' capsule the one with the maximum objective function value when scored.

Reinforcement Learning

Rules Engine

`jstate diagram`

Parameter Tuning

Other Methods

Expectimax

tree would not be deep to merit this.

Alpha-Beta Search

Russell & Norvig, pg 166

Hidden Markov Model

Roland Memisevic <http://www.iro.umontreal.ca/~memisevr/code/hmm.py>

Needed to find a HMM implementation that respected the dependencies of the game cluster.
(For example could not use GHMM www.ghmm.org due to a C library dependency)

Conclusion

A combination of ML and traditional techniques is more powerful than using one set of techniques alone.

References

- [1] *Reinforcement Learning*, Sutton & Barto, 1998, ISBN-10: 0-262-19398-1