

CSCI E-181 Spring 2014 Practical 1

David Wihl
davidwihl@gmail.com

February 9, 2014

Warm-Up

As a warmup, I synthesized five clusters of data. I then used a K-Means implementation in Octave I had written for a previous course.¹ While this implementation was sufficient for the prior course's Dataset, when I tested it the synthesized data set, $K=5$ and random initial centroids, one of the centroids would frequently not converge on any points.




Figure 1: Random Initial Centroids After 1 Iteration

¹Machine Learning, Coursera, Prof. Andrew Ng, Completed Jan 2014, <https://class.coursera.org/ml-004>

I subsequently modified the code to use K-Medoids, choosing one of the sample data points at random as an initial centroid. This worked much better.




Figure 2: K-Medoids Converge After 4 Iterations

CIFAR-10 Image Data

I then attempted using K-Medoids with the CIFAR-10 Image Data, using the Matlab version of the data with Octave. The training data consists of a 10000x3072 matrix of UInt8. Each row is a 32x32x3 (=3072) color image, consisting of 1024 red, 1024 green and 1024 blue elements. There are 10 classes in the set (“airplane”, “automobile”, etc.), so setting K=10 was a rational first step.

Percentage Distribution of K values after normalization and 10 iterations

06 05 04 26 14 13 05 04 03 15